

CONFIGURATION MANAGEMENT SYSTEM

The configuration management system exists to provide every daemon with the proper configuration information. The configuration can be viewed as a set of key-value pairs.

How can the configuration be set? Well, there are several sources:

- the ceph configuration file, usually named `ceph.conf`
- command line arguments::
 `-debug-ms=1 -debug-pg=10` etc.
- arguments injected at runtime by using `injectargs`

THE CONFIGURATION FILE

Most configuration settings originate in the Ceph configuration file.

How do we find the configuration file? Well, in order, we check:

- the default locations
- the environment variable `CEPH_CONF`
- the command line argument `-c`

Each stanza of the configuration file describes the key-value pairs that will be in effect for a particular subset of the daemons. The “global” stanza applies to everything. The “mon”, “osd”, and “mds” stanzas specify settings to take effect for all monitors, all osds, and all mds servers, respectively. A stanza of the form `mon.$name`, `osd.$name`, or `mds.$name` gives settings for the monitor, OSD, or MDS of that name, respectively. Configuration values that appear later in the file win over earlier ones.

A sample configuration file can be found in `src/sample.ceph.conf`.

METAVARIABLES

The configuration system allows any configuration value to be substituted into another value using the `$varname` syntax, similar to how bash shell expansion works.

A few additional special metavariables are also defined:

- `$host`: expands to the current hostname
- `$type`: expands to one of “mds”, “osd”, “mon”, or “client”
- `$id`: expands to the daemon identifier. For `osd.0`, this would be `0`; for `mds.a`, it would be `a`; for `client.admin`, it would be `admin`.
- `$num`: same as `$id`
- `$name`: expands to `$type.$id`

READING CONFIGURATION VALUES

There are two ways for Ceph code to get configuration values. One way is to read it directly from a variable named “`g_conf`,” or equivalently, “`g_ceph_ctx->_conf`.” The other is to register an observer that will be called every time the relevant configuration values change. This observer will be called soon after the initial configuration is read, and every time after that when one of the relevant values changes. Each observer tracks a set of keys and is invoked only when one of the relevant keys changes.

The interface to implement is found in `common/config_obs.h`.

The observer method should be preferred in new code because

- It is more flexible, allowing the code to do whatever reinitialization needs to be done to implement the new configuration value.
- It is the only way to create a `std::string` configuration variable that can be changed by `injectargs`.
- Even for int-valued configuration options, changing the values in one thread while another thread is reading them can lead to subtle and impossible-to-diagnose bugs.

For these reasons, reading directly from `g_conf` should be considered deprecated and not done in new code. Do not ever alter `g_conf`.

CHANGING CONFIGURATION VALUES

Configuration values can be changed by calling `g_conf->set_val`. After changing the configuration, you should call `g_conf->apply_changes` to re-run all the affected configuration observers. For convenience, you can call `g_conf->set_val_or_die` to make a configuration change which you think should never fail.

`Injectargs`, `parse_argv`, and `parse_env` are three other functions which modify the configuration. Just like with `set_val`, you should call `apply_changes` after calling these functions to make sure your changes get applied.