

## CONTROL COMMANDS

### MONITOR COMMANDS

Monitor commands are issued using the ceph utility:

```
ceph [-m monhost] {command}
```

The command is usually (though not always) of the form:

```
ceph {subsystem} {command}
```

### SYSTEM COMMANDS

Execute the following to display the current status of the cluster.

```
ceph -s  
ceph status
```

Execute the following to display a running summary of the status of the cluster, and major events.

```
ceph -w
```

Execute the following to show the monitor quorum, including which monitors are participating and which one is the leader.

```
ceph quorum_status
```

Execute the following to query the status of a single monitor, including whether or not it is in the quorum.

```
ceph [-m monhost] mon_status
```

### AUTHENTICATION SUBSYSTEM

To add a keyring for an OSD, execute the following:

```
ceph auth add {osd} [--in-file|-i] {path-to-osd-keyring}
```

To list the cluster's keys and their capabilities, execute the following:

```
ceph auth list
```

### PLACEMENT GROUP SUBSYSTEM

To display the statistics for all placement groups, execute the following:

```
ceph pg dump [--format {format}]
```

The valid formats are plain (default) and json.

To display the statistics for all placement groups stuck in a specified state, execute the following:

```
ceph pg dump_stuck inactive|unclean|stale [--format {format}] [-t|--threshold {seconds}]
```

--format may be plain (default) or json

--threshold defines how many seconds “stuck” is (default: 300)

**Inactive** Placement groups cannot process reads or writes because they are waiting for an OSD with the most up-to-date data to come back.

**Unclean** Placement groups contain objects that are not replicated the desired number of times. They should be recovering.

**Stale** Placement groups are in an unknown state - the OSDs that host them have not reported to the monitor cluster in a while (configured by mon\_osd\_report\_timeout).

Revert “lost” objects to their prior state, either a previous version or delete them if they were just created.

```
ceph pg {pgid} mark_unfound_lost revert
```

## OSD SUBSYSTEM

Query osd subsystem status.

```
ceph osd stat
```

Write a copy of the most recent osd map to a file. See [osdmaptool](#).

```
ceph osd getmap -o file
```

Write a copy of the crush map from the most recent osd map to file.

```
ceph osd getcrushmap -o file
```

The foregoing functionally equivalent to

```
ceph osd getmap -o /tmp/osdmap  
osdmaptool /tmp/osdmap --export-crush file
```

Dump the OSD map. Valid formats for -f are plain and json. If no --format option is given, the OSD map is dumped as plain text.

```
ceph osd dump [--format {format}]
```

Dump the OSD map as a tree with one line per OSD containing weight and state.

```
ceph osd tree [--format {format}]
```

Find out where a specific object is or would be stored in the system:

```
ceph osd map <pool-name> <object-name>
```

Add or move a new item (OSD) with the given id/name/weight at the specified location.

```
ceph osd crush set {id} {weight} [{loc1}] [{loc2} ...]
```

Remove an existing item from the CRUSH map.

```
ceph osd crush remove {id}
```

Move an existing bucket from one position in the hierarchy to another.

```
ceph osd crush move {id} {loc1} [{loc2} ...]
```

Set the weight of the item given by {name} to {weight}.

```
ceph osd crush reweight {name} {weight}
```

Create a cluster snapshot.

```
ceph osd cluster_snap {name}
```

Mark an OSD as lost. This may result in permanent data loss. Use with caution.

```
ceph osd lost {id} [--yes-i-really-mean-it]
```

Create a new OSD. If no UUID is given, it will be set automatically when the OSD starts up.

```
ceph osd create [{uuid}]
```

Remove the given OSD(s).

```
ceph osd rm [{id}...]
```

Query the current max\_osd parameter in the osd map.

```
ceph osd getmaxosd
```

Import the given OSD map. Note that this can be a bit dangerous, since the OSD map includes dynamic state about which OSDs are current on or offline; only do this if you've just modified a (very) recent copy of the map.

```
ceph osd setmap -i file
```

Import the given crush map.

```
ceph osd setcrushmap -i file
```

Set the max\_osd parameter in the OSD map. This is necessary when expanding the storage cluster.

```
ceph osd setmaxosd
```

Mark OSD {osd-num} down.

```
ceph osd down {osd-num}
```

Mark OSD {osd-num} out of the distribution (i.e. allocated no data).

---

```
ceph osd out {osd-num}
```

Mark {osd-num} in the distribution (i.e. allocated data).

```
ceph osd in {osd-num}
```

List classes that are loaded in the ceph cluster.

```
ceph class list
```

Set or clear the pause flags in the OSD map. If set, no IO requests will be sent to any OSD. Clearing the flags via unpause results in resending pending requests.

```
ceph osd pause  
ceph osd unpause
```

Set the weight of {osd-num} to {weight}. Two OSDs with the same weight will receive roughly the same number of I/O requests and store approximately the same amount of data.

```
ceph osd reweight {osd-num} {weight}
```

Reweights all the OSDs by reducing the weight of OSDs which are heavily overused. By default it will adjust the weights downward on OSDs which have 120% of the average utilization, but if you include threshold it will use that percentage instead.

```
ceph osd reweight-by-utilization [threshold]
```

Adds/removes the address to/from the blacklist. When adding an address, you can specify how long it should be blacklisted in seconds; otherwise, it will default to 1 hour. A blacklisted address is prevented from connecting to any OSD. Blacklisting is most often used to prevent a lagging metadata server from making bad changes to data on the OSDs.

These commands are mostly only useful for failure testing, as blacklists are normally maintained automatically and shouldn't need manual intervention.

```
ceph osd blacklist add ADDRESS[:source_port] [TIME]  
ceph osd blacklist rm ADDRESS[:source_port]
```

Creates/deletes a snapshot of a pool.

```
ceph osd pool mksnap {pool-name} {snap-name}  
ceph osd pool rmsnap {pool-name} {snap-name}
```

Creates/deletes/renames a storage pool.

```
ceph osd pool create {pool-name} pg_num [pgp_num]  
ceph osd pool delete {pool-name} [{pool-name} --yes-i-really-really-mean-it]  
ceph osd pool rename {old-name} {new-name}
```

Changes a pool setting.

```
ceph osd pool set {pool-name} {field} {value}
```

Valid fields are:

- `size`: Sets the number of copies of data in the pool.
- `crashreplay_interval`: The number of seconds to allow clients to replay acknowledged but uncommitted

requests.

- pg\_num: The placement group number.
- pgp\_num: Effective number when calculating pg placement.
- crush\_ruleset: rule number for mapping placement.

Get the value of a pool setting.

```
ceph osd pool get {pool-name} {field}
```

Valid fields are:

- pg\_num: The placement group number.
- pgp\_num: Effective number of placement groups when calculating placement.
- lpg\_num: The number of local placement groups.
- lpgp\_num: The number used for placing the local placement groups.

Sends a scrub command to OSD {osd-num}. To send the command to all OSDs, use \*.

```
ceph osd scrub {osd-num}
```

Sends a repair command to osdN. To send the command to all osds, use \*.

```
ceph osd repair N
```

Runs a simple throughput benchmark against osdN, writing TOTAL\_BYTES in write requests of BYTES\_PER\_WRITE each. By default, the test writes 1 GB in total in 4-MB increments.

```
ceph osd tell N bench [BYTES_PER_WRITE] [TOTAL_BYTES]
```

## MDS SUBSYSTEM

Change configuration parameters on a running mds.

```
ceph mds tell {mds-id} injectargs '--{switch} {value} [--{switch} {value}]'
```

Example:

```
ceph mds tell 0 injectargs '--debug_ms 1 --debug_mds 10'
```

Enables debug messages.

```
ceph mds stat
```

Displays the status of all metadata servers.

```
ceph mds fail 0
```

Marks the active MDS as failed, triggering failover to a standby if present.

**Todo:** ceph mds subcommands missing docs: set\_max\_mds, dump, getmap, stop, setmap

## MON SUBSYSTEM

Show monitor stats:

```
ceph mon stat
```

```
2011-12-14 10:40:59.044395 mon {- [mon,stat]
2011-12-14 10:40:59.057111 mon.1 -} 'e3: 5 mons at {a=10.1.2.3:6789/0,b=10.1.2.4:6789/0,c=10.
```

The quorum list at the end lists monitor nodes that are part of the current quorum.

This is also available more directly:

```
$ ./ceph quorum_status
```

```
2011-12-14 10:44:20.417705 mon {- [quorum_status]
2011-12-14 10:44:20.431890 mon.0 -}
```

```
{ "election_epoch": 10,
  "quorum": [
    0,
    1,
    2],
  "monmap": { "epoch": 1,
    "fsid": "444b489c-4f16-4b75-83f0-cb8097468898",
    "modified": "2011-12-12 13:28:27.505520",
    "created": "2011-12-12 13:28:27.505520",
    "mons": [
      { "rank": 0,
        "name": "a",
        "addr": "127.0.0.1:6789\0"},
      { "rank": 1,
        "name": "b",
        "addr": "127.0.0.1:6790\0"},
      { "rank": 2,
        "name": "c",
        "addr": "127.0.0.1:6791\0"}]}}' (0)
```

The above will block until a quorum is reached.

For a status of just the monitor you connect to (use -m HOST:PORT to select):

```
ceph mon_status
```

```
2011-12-14 10:45:30.644414 mon {- [mon_status]
2011-12-14 10:45:30.644632 mon.0 -}
```

```
{ "name": "a",
  "rank": 0,
  "state": "leader",
  "election_epoch": 10,
  "quorum": [
    0,
    1,
    2],
  "outside_quorum": [],
  "monmap": { "epoch": 1,
    "fsid": "444b489c-4f16-4b75-83f0-cb8097468898",
    "modified": "2011-12-12 13:28:27.505520",
    "created": "2011-12-12 13:28:27.505520",
    "mons": [
      { "rank": 0,
        "name": "a",
        "addr": "127.0.0.1:6789\0"},
      { "rank": 1,
        "name": "b",
        "addr": "127.0.0.1:6790\0"},
      { "rank": 2,
        "name": "c",
        "addr": "127.0.0.1:6791\0"}]}}' (0)
```

A dump of the monitor state:

```
ceph mon dump
```

```
2011-12-14 10:43:08.015333 mon {- [mon,dump]
2011-12-14 10:43:08.015567 mon.0 -} 'dumped monmap epoch 1' (0)
epoch 1
fsid 444b489c-4f16-4b75-83f0-cb8097468898
last_changed 2011-12-12 13:28:27.505520
created 2011-12-12 13:28:27.505520
0: 127.0.0.1:6789/0 mon.a
1: 127.0.0.1:6790/0 mon.b
2: 127.0.0.1:6791/0 mon.c
```