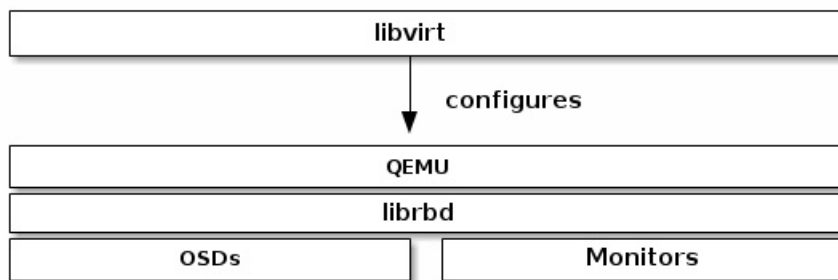


USING LIBVIRT WITH CEPH RBD

The libvirt library creates a virtual machine abstraction layer between hypervisor interfaces and the software applications that use them. With libvirt, developers and system administrators can focus on a common management framework, common API, and common shell interface (i.e., virsh) to many different hypervisors, including:

- QEMU/KVM
- XEN
- LXC
- VirtualBox
- etc.

Ceph block devices support QEMU/KVM. You can use Ceph block devices with software that interfaces with libvirt. The following stack diagram illustrates how libvirt and QEMU use Ceph block devices via librbd.



The most common libvirt use case involves providing Ceph block devices to cloud solutions like OpenStack or CloudStack. The cloud solution uses libvirt to interact with QEMU/KVM, and QEMU/KVM interacts with Ceph block devices via librbd. See [Block Devices and OpenStack](#) and [Block Devices and CloudStack](#) for details.

You can also use Ceph block devices with libvirt, virsh and the libvirt API. See [libvirt Virtualization API](#) for details.

PREREQUISITES

- [Install](#) and [configure](#) a Ceph cluster
- [Install](#) and [configure](#) QEMU/KVM

INSTALLING LIBVIRT ON UBUNTU 12.04 PRECISE

libvirt packages are incorporated into the Ubuntu 12.04 precise distribution. To install libvirt on precise, execute the following:

```
sudo apt-get update && sudo apt-get install libvirt-bin
```

INSTALLING LIBVIRT ON EARLIER VERSIONS OF UBUNTU

For Ubuntu distributions 11.10 oneiric and earlier, you must build libvirt from source. Clone the libvirt repository, and use [AutoGen](#) to generate the build. Then, execute `make` and `make install` to complete the installation. For example:

```
git clone git://libvirt.org/libvirt.git
cd libvirt
./autogen.sh
make
sudo make install
```

See [libvirt Installation](#) for details.

USING CEPH WITH VIRTUAL MACHINES

To create VMs that use Ceph block devices, use the procedures in the following sections. In the exemplary embodiment, we've used `libvirt-pool` for the pool name, `client.libvirt` for the user name, and `new-libvirt-image` for the image name. You may use any value you like, but ensure you replace those values when executing commands in the subsequent procedures.

CONFIGURING CEPH

To configure Ceph for use with libvirt, perform the following steps:

1. **Create a pool** (or use the default). The following example uses the pool name `libvirt-pool` with 128 placement groups.

```
ceph osd pool create libvirt-pool 128 128
```

Verify the pool exists.

```
ceph osd lspools
```

2. **Create a Ceph Name** (or use `client.admin` for version 0.9.7 and earlier). The following example uses the Ceph name `client.libvirt` and references `libvirt-pool`.

```
ceph auth get-or-create client.libvirt mon 'allow r' osd 'allow class-read object_prefix
```

Verify the name exists.

```
ceph auth list
```

NOTE: libvirt will access Ceph using the ID `libvirt`, not the Ceph name `client.libvirt`. See [Ceph Commandline](#) for detailed explanation of the difference between ID and name.

3. Use QEMU to **create an image** in your RBD pool. The following example uses the image name `new-libvirt-image` and references `libvirt-pool`.

```
qemu-img create -f rbd rbd:libvirt-pool/new-libvirt-image 2G
```

Verify the image exists.

```
rbd -p libvirt-pool ls
```

NOTE: You can also use `rbd create` to create an image, but we recommend ensuring that QEMU is working properly.

PREPARING THE VM MANAGER

You may use libvirt without a VM manager, but you may find it simpler to create your first domain with `virt-manager`.

1. Install a virtual machine manager. See [KVM/VirtManager](#) for details.

```
sudo apt-get install virt-manager
```

2. Download an OS image (if necessary).
3. Launch the virtual machine manager.

```
sudo virt-manager
```

CREATING A VM

To create a VM with `virt-manager`, perform the following steps:

1. Press the **Create New Virtual Machine** button.
2. Name the new virtual machine domain. In the exemplary embodiment, we use the name `libvirt-virtual-machine`. You may use any name you wish, but ensure you replace `libvirt-virtual-machine` with the name you choose in subsequent commandline and configuration examples.

```
libvirt-virtual-machine
```

3. Import the image.

```
/path/to/image/recent-linux.img
```

NOTE: Import a recent image. Some older images may not rescan for virtual devices properly.

4. Configure and start the VM.
5. You may use `virsh list` to verify the VM domain exists.

```
sudo virsh list
```

6. Login to the VM (root/root)
7. Stop the VM before configuring it for use with Ceph.

CONFIGURING THE VM

When configuring the VM for use with Ceph, it is important to use `virsh` where appropriate. Additionally, `virsh` commands often require root privileges (i.e., `sudo`) and will not return appropriate results or notify you that that root privileges are required. For a reference of `virsh` commands, refer to [Virsh Command Reference](#).

1. Open the configuration file with `virsh edit`.

```
sudo virsh edit {vm-domain-name}
```

Under `<devices>` there should be a `<disk>` entry.

```
<devices>
  <emulator>/usr/bin/kvm</emulator>
  <disk type='file' device='disk'>
    <driver name='qemu' type='raw'/>
    <source file='/path/to/image/recent-linux.img'/>
    <target dev='vda' bus='virtio'/>
    <address type='drive' controller='0' bus='0' unit='0'/>
  </disk>
```

Replace `/path/to/image/recent-linux.img` with the path to the OS image. The minimum kernel for using the faster `virtio` bus is 2.6.25. See [Virtio](#) for details.

IMPORTANT: Use `sudo virsh edit` instead of a text editor. If you edit the configuration file under `/etc/libvirt/qemu` with a text editor, `libvirt` may not recognize the change. If there is a discrepancy between the contents of the XML file under `/etc/libvirt/qemu` and the result of `sudo virsh dumpxml {vm-domain-name}`, then your VM may not work properly.

2. Add the Ceph RBD image you created as a `<disk>` entry.

```
<disk type='network' device='disk'>
  <source protocol='rbd' name='libvirt-pool/new-libvirt-image'>
    <host name='{monitor-host}' port='6789'/>
```

```
        </source>
        <target dev='hdb' bus='ide' />
</disk>
```

Replace {monitor-host} with the name of your host, and replace the pool and/or image name as necessary. You may add multiple <host> entries for your Ceph monitors. The dev attribute is the logical device name that will appear under the /dev directory of your VM. The optional bus attribute indicates the type of disk device to emulate. The valid settings are driver specific (e.g., "ide", "scsi", "virtio", "xen", "usb" or "sata").

See [Disks](#) for details of the <disk> element, and its child elements and attributes.

3. Save the file.
4. If you are using [Ceph Authentication](#), you must generate a secret.

```
cat > secret.xml <<EOF
<secret ephemeral='no' private='no'>
  <usage type='ceph'>
    <name>client.libvirt secret</name>
  </usage>
</secret>
EOF
```

5. Define the secret.

```
sudo virsh secret-define --file secret.xml
<uuid of secret is output here>
```

6. Get the client.libvirt key and save the key string to a file.

```
sudo ceph auth list
vim client.libvirt.key
```

7. Set the UUID of the secret.

```
sudo virsh secret-set-value --secret {uuid of secret} --base64 $(cat client.libvirt.key)
```

You must also set the secret manually by adding the following <auth> entry to the <disk> element you entered earlier (replacing the uuid value with the result from the command line example above).

```
sudo virsh edit {vm-domain-name}
```

Then, add <auth></auth> element to the domain configuration file:

```
...
</source>
<auth username='libvirt'>
  <secret type='ceph' uuid='9ec59067-fdbc-a6c0-03ff-df165c0587b8' />
</auth>
<target ...
```

NOTE: The exemplary ID is libvirt, not the Ceph name client.libvirt as generated at step 2 of [Configuring Ceph](#). Ensure you use the ID component of the Ceph name you generated. If for some reason you need to regenerate the secret, you will have to execute `sudo virsh secret-undefine {uuid}` before executing `sudo virsh secret-set-value` again.

SUMMARY

Once you have configured the VM for use with Ceph, you can start the VM. To verify that the VM and Ceph are communicating,

you may perform the following procedures.

1. Check to see if Ceph is running:

```
ceph health
```

2. Check to see if the VM is running.

```
sudo virsh list
```

3. Check to see if the VM is communicating with Ceph. Replace {vm-domain-name} with the name of your VM domain:

```
sudo virsh qemu-monitor-command --hmp {vm-domain-name} 'info block'
```

4. Check to see if the device from `<target dev='hdb' bus='ide' />` appears under `/dev` or under `proc/partitions`.

```
ls dev  
cat proc/partitions
```

If everything looks okay, you may begin using the Ceph block device within your VM.
