

INTERNAL DEVELOPER DOCUMENTATION

Note: If you're looking for how to use Ceph as a library from your own software, please see [API Documentation](#).

You can start a development mode Ceph cluster, after compiling the source, with:

```
cd src
install -d -m0755 out dev/osd0
./vstart.sh -n -x -l
# check that it's there
./ceph health
```

Todo: vstart is woefully undocumented and full of sharp sticks to poke yourself with.

Mailing list

The official development email list is `ceph-devel@vger.kernel.org`. Subscribe by sending a message to `majordomo@vger.kernel.org` with the line:

```
subscribe ceph-devel
```

in the body of the message.

Contents

- [Tracing Ceph With BlkKin](#)
 - [Installing Blkin](#)
 - [Configuring Ceph with Blkin](#)
 - [Testing Blkin](#)
 - [Install Zipkin](#)
 - [Show Ceph's Blkin Traces in Zipkin-web](#)
- [BlueStore Internals](#)
 - [Small write strategies](#)
 - [Possible future modes](#)
 - [Mapping](#)
- [Cache pool](#)
 - [Purpose](#)
 - [Use cases](#)
 - [Read-write pool, writeback](#)
 - [Read-only pool, weak consistency](#)
 - [Interface](#)
 - [Tiering agent](#)
 - [HitSet metadata](#)
 - [Cache mode](#)
 - [Cache sizing](#)
 - [Other tunables](#)
- [ceph-disk](#)
 - [device-mapper crypt](#)
 - [Settings](#)
 - [lockbox](#)
 - [partitions](#)
 - [prepare class hierarchy](#)
 - [state transition of partitions](#)
- [CephFS Snapshots](#)
 - [Overview](#)
 - [Important Data Structures](#)
 - [Creating a snapshot](#)
 - [Updating a snapshot](#)
 - [Generating a SnapContext](#)
 - [Storing snapshot data](#)

- Storing snapshot metadata
 - Snapshot writeback
 - Deleting snapshots
 - Hard links
 - Multi-FS
- A Detailed Description of the Cephx Authentication Protocol
 - Introduction
 - Getting Started With Authorization
 - Phase I:
 - Phase II
- Configuration Management System
 - The Configuration File
 - Metavariables
 - Reading configuration values
 - Changing configuration values
 - Defining config options
 - Levels
 - Description and long description
 - Default values
 - Safety
 - Service
 - Tags
 - Enums
- config-key layout
 - OSD dm-crypt keys
 - ceph-mgr modules
 - Configuration
- CephContext
- Corpus structure
 - How to generate an object corpus
- Installing Oprofile
 - Installation
 - Compiling Ceph for Profiling
 - Ceph Configuration
- C++17 and libstdc++ ABI
- CephFS delayed deletion
- Deploying a development cluster
 - Usage
 - Options
 - Environment variables
- Deploying multiple development clusters on the same machine
 - Usage
- Development workflows
 - Release Cycle
 - Merging bug fixes or features
 - Resolving bug reports and implementing features
 - Running and interpreting teuthology integration tests
 - Preparing a new release
 - Cutting a new stable release
 - Cutting a new development release
 - Publishing point releases and backporting
- Documenting Ceph
 - User documentation
 - Code Documentation
 - Drawing diagrams
 - Graphviz
 - Dita
 - Blockdiag
 - Inkscape
- Serialization (encode/decode)
 - Adding a field to a structure
- Erasure Coded pool
 - Purpose
 - Use cases
 - Cold storage

- Cheap multidatacenter storage
- Interface
- File striping
 - ceph_file_layout
- FreeBSD Implementation details
 - Disk layout
 - Configuration
 - MON creation
 - OSD creation
- Building Ceph Documentation
 - Clone the Ceph Repository
 - Install the Required Tools
 - Build the Documents
- IANA Numbers
 - Private Enterprise Number (PEN) Assignment
 - Port number (monitor)
- Contributing to Ceph: A Guide for Developers
 - Introduction
 - Essentials (tl;dr)
 - Leads
 - History
 - Licensing
 - Source code repositories
 - Redmine issue tracker
 - Mailing list
 - IRC
 - Submitting patches
 - Building from source
 - Using ccache to speed up local builds
 - Development-mode cluster
 - Backporting
 - Guidance for use of cluster log
 - What is merged where and when ?
 - Development releases (i.e. x.0.z)
 - What ?
 - Where ?
 - When ?
 - Branch merges
 - Stable release candidates (i.e. x.1.z) phase 1
 - What ?
 - Where ?
 - When ?
 - Branch merges
 - Stable release candidates (i.e. x.1.z) phase 2
 - What ?
 - Where ?
 - When ?
 - Branch merges
 - Stable releases (i.e. x.2.z)
 - What ?
 - Where ?
 - When ?
 - Branch merges
 - Issue tracker
 - Issue tracker conventions
 - Basic workflow
 - Update the tracker
 - Upstream code
 - Local environment
 - Bugfix branch
 - Fix bug locally
 - GitHub pull request
 - Automated PR validation
 - Notes on PR make check test
 - Integration tests AKA ceph-qa-suite

- Code review
 - Amending your PR
 - Merge
- Testing
- Testing - make check
 - Caveats
- Testing - integration tests
 - Teuthology consumes packages
 - The nightlies
 - Suites inventory
 - teuthology-describe-tests
 - How integration tests are run
 - How integration tests are defined
 - Reading a standalone test
 - Test descriptions
 - How are tests built from directories?
 - Convolution operator
 - Concatenation operator
 - Filtering tests by their description
 - Reducing the number of tests
- Testing in the cloud
 - Assumptions and caveat
 - Prepare tenant
 - Getting ceph-workbench
 - Linking ceph-workbench with your OpenStack tenant
 - Run the dummy suite
 - Run a standalone test
 - Interrupt a running suite
 - Upload logs to archive server
 - Provision VMs ad hoc
 - Deploy a cluster for manual testing
- Testing - how to run s3-tests locally
 - Step 1 - build Ceph
 - Step 2 - vstart
 - Step 3 - run s3-tests
- Kernel client troubleshooting (FS)
- Library architecture
- Use of the cluster log
 - Severity
 - Frequency
 - Language and formatting
- Debug logs
 - Performance counters
- build on MacOS
- Messenger notes
 - ceph_perf_msgr
- Monitor bootstrap
 - Logical id
 - Secret keys
 - Cluster fsid
 - Monitor address
 - Peers
 - Cluster creation
 - Names and addresses
 - Addresses only
 - Names only
 - Cluster expansion
 - Initially peerless expansion
 - Expanding with initial members
- msgr2 protocol
 - Goals
 - Definitions
 - Phases
 - Banner
 - Frame format

- Authentication
 - Message frame format
 - Message flow handshake
 - Message exchange
- Network Encoding
 - Conventions
 - Integers
 - Complex Types
 - Variable Arrays
 - Primitive Aliases
 - Structures
 - Optional
 - Pair
 - Triple
 - List
 - Blob
 - Map
 - Complex Types
 - `utime_t`
 - `ceph_entity_name`
- Network Protocol
 - Hello
 - Banner
 - Connect
 - Connect Reply
 - MSGR Protocol
 - `CEPH_MSGR_TAG_CLOSE` (0x06)
 - `CEPH_MSGR_TAG_MSG` (0x07)
 - `CEPH_MSGR_TAG_ACK` (0x08)
 - `CEPH_MSGR_TAG_KEEPALIVE` (0x09)
 - `CEPH_MSGR_TAG_KEEPALIVE2` (0x0E)
 - `CEPH_MSGR_TAG_KEEPALIVE2_ACK` (0x0F)
- Object Store Architecture Overview
- OSD class path issues
- Peering
 - Concepts
 - Description of the Peering Process
 - State Model
- Using perf
 - Common Issues
 - Flamegraphs
- Perf counters
 - Access
 - Collections
 - Schema
 - Dump
- Perf histograms
 - Access
 - Collections
 - Schema
 - Dump
- PG (Placement Group) notes
 - Overview
 - Mapping algorithm (simplified)
 - User-visible PG States
- Developer Guide (Quick)
 - Development
 - Running a development deployment
 - Resetting your vstart environment
 - Running a RadosGW development environment
 - Run unit tests
- RADOS client protocol
 - Basics
 - Resends
 - Backoff

- RBD Incremental Backup
 - Header
 - Metadata records
 - From snap
 - To snap
 - Size
 - Data Records
 - Updated data
 - Zero data
 - Final Record
 - End
 - Header
 - Metadata records
 - From snap
 - To snap
 - Size
 - Data Records
 - Updated data
 - Zero data
 - Final Record
 - End
- RBD Export & Import
 - Header
 - Metadata records
 - Image order
 - Image format
 - Image Features
 - Image Stripe unit
 - Image Stripe count
 - ImageMeta Key and Value
 - Final Record
 - End
 - Diffs records
 - Snap Protection status
 - Others
- RBD Layering
 - Command line interface
 - Implementation
 - Data Flow
 - Parent/Child relationships
 - Protection
 - Resizing
 - Renaming
 - Header changes
 - cls_rbd
 - librbd
- Ceph Release Process
 - 1. Build environment
 - 2. Setup keyring for signing packages
 - 3. Set up build area
 - 4. Update Build version numbers
 - 5. Create Makefiles
 - 6. Run the release scripts
 - 7. Create RPM Repo
 - 8. Create Debian repo
 - 9. Push repos to ceph.org
 - 10. Update Git
 - Point release
 - Development and Stable releases
- Notes on Ceph repositories
 - Special branches
 - Rules
- Sepia community test lab
- Session Authentication for the Cephx Protocol
 - Introduction

- [Storing the Key](#)
- [Signing Messages](#)
- [Checking Signatures](#)
- [Adding New Session Authentication Methods](#)
- [Adding Encryption to Sessions](#)
- [Session Security Statistics](#)
- [Testing notes](#)
 - [build-integration-branch](#)
 - [Setup](#)
 - [Using](#)
- [Public OSD Version](#)
- [Wireshark Dissector](#)
 - [Using](#)
 - [Developing](#)
- [OSD developer documentation](#)
 - [Backfill Reservation](#)
 - [Erasure Coded Placement Groups](#)
 - [Glossary](#)
 - [Table of content](#)
 - [Developer notes](#)
 - [Introduction](#)
 - [Reading and writing encoded chunks from and to OSDs](#)
 - [Erasure code library](#)
 - [Notes](#)
 - [Jerasure plugin](#)
 - [Introduction](#)
 - [ECBackend Implementation Strategy](#)
 - [Misc initial design notes](#)
 - [PGTemp and Crush](#)
 - [Distinguished acting set positions](#)
 - [Object Classes](#)
 - [Scrub](#)
 - [Crush](#)
 - [ECBackend](#)
 - [MAIN OPERATION OVERVIEW](#)
 - [WHOLE STRIPE WRITE](#)
 - [READ-MODIFY-WRITE](#)
 - [OSD Object Write and Consistency](#)
 - [ExtentCache](#)
 - [Pipeline](#)
- [last_epoch_started](#)
- [Log Based PG](#)
 - [Background](#)
 - [Why PrimaryLogPG?](#)
 - [Primary log-based replication](#)
 - [ReplicatedBackend/ECBackend unification strategy](#)
 - [PGBackend](#)
 - [PGBackend Interface Explanation](#)
 - [Readable vs Degraded](#)
 - [Client Reads](#)
 - [Scrub](#)
 - [Recovery](#)
- [Map and PG Message handling](#)
 - [Overview](#)
 - [MOSDMap](#)
 - [MOSDPGOp/MOSDPGSubOp](#)
 - [CEPH_MSG_OSD_OP processing](#)
 - [Peering Messages](#)
- [OSD](#)
 - [Concepts](#)
 - [Overview](#)
- [OSD Throttles](#)
 - [WBThrottle](#)
 - [op_queue_throttle](#)
 - [journal usage throttle](#)

- PG
 - Concepts
 - Peering Details and Gotchas
- PG Removal
- PGPool
- Recovery Reservation
 - Things to Note
 - See Also
- Scrubbing Behavior Table
 - State variables
- Snaps
 - Overview
 - Ondisk Structures
 - Snap Removal
 - Recovery
 - SnapMapper
 - Split
- Watch Notify
 - Overview
 - Watch Lifecycle
 - Notify Lifecycle
- Writeback Throttle
- MDS developer documentation
 - MDS internal data structures
 - Subtree exports
 - Normal Migration
- RADOS Gateway developer documentation
 - Usage Design Overview
 - Testing
 - Admin Ops Nonimplemented
 - Get Object
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
 - Head Object
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
 - Get Zone Info
 - Syntax
 - Response Entities
 - Special Error Responses
 - Example Response
 - Add Placement Pool
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
 - Remove Placement Pool
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
 - List Available Data Placement Pools
 - Syntax
 - Response Entities
 - List Expired Garbage Collection Items
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
 - Manually Processes Garbage Collection Items
 - Syntax

- Request Parameters
 - Response Entities
 - Special Error Responses
- Show Log Objects
 - Syntax
 - Request Parameters
 - Response Entities
 - Special Error Responses
- Standard Error Responses
- Rados Gateway S3 API Compliance
 - Naming code reference
 - Common Request Headers
 - Common Response Headers
 - Operations on the Service
 - Operations on Buckets
 - Operations on Objects
- ceph-volume developer documentation
 - Plugins
 - Adding Plugins
 - LVM
 - Tag API
 - Metadata
 - **type**
 - **cluster_fsid**
 - **data_device**
 - **data_uuid**
 - **journal_device**
 - **journal_uuid**
 - **encrypted**
 - **osd_fsid**
 - **osd_id**
 - **block_device**
 - **block_uuid**
 - **db_device**
 - **db_uuid**
 - **wal_device**
 - **wal_uuid**
 - systemd