

PERF HISTOGRAMS

The perf histograms build on perf counters infrastructure. Histograms are built for a number of counters and simplify gathering data on which groups of counter values occur most often over time. Perf histograms are currently unsigned 64-bit integer counters, so they're mostly useful for time and sizes. Data dumped by perf histogram can then be feed into other analysis tools/scripts.

ACCESS

The perf histogram data are accessed via the admin socket. For example:

```
ceph daemon osd.0 perf histogram schema
ceph daemon osd.0 perf histogram dump
```

COLLECTIONS

The histograms are grouped into named collections, normally representing a subsystem or an instance of a subsystem. For example, the internal throttle mechanism reports statistics on how it is throttling, and each instance is named something like:

```
op_r_latency_out_bytes_histogram
op_rw_latency_in_bytes_histogram
op_rw_latency_out_bytes_histogram
...
```

SCHEMA

The `perf histogram schema` command dumps a json description of which values are available, and what their type is. Each named value as a type bitfield, with the 5-th bit always set and following bits defined.

| bit | meaning |
|-----|---------|
|-----|---------|

| | |
|---|-------------------------------|
| 1 | floating point value |
| 2 | unsigned 64-bit integer value |
| 4 | average (sum + count pair) |
| 8 | counter (vs gauge) |

In other words, histogram of type “18” is a histogram of unsigned 64-bit integer values (16 + 2).

Here is an example of the schema output:

```
{
  "AsyncMessenger::Worker-0": {},
  "AsyncMessenger::Worker-1": {},
  "AsyncMessenger::Worker-2": {},
  "mutex-WBThrottle::lock": {},
  "objecter": {},
  "osd": {
    "op_r_latency_out_bytes_histogram": {
      "type": 18,
      "description": "Histogram of operation latency (including queue time) + da ta",
      "nick": ""
    },
    "op_w_latency_in_bytes_histogram": {
      "type": 18,
      "description": "Histogram of operation latency (including queue time) + da ta",
      "nick": ""
    },
    "op_rw_latency_in_bytes_histogram": {
      "type": 18,
      "description": "Histogram of rw operation latency (including queue time) + da
```

```

        "nick": ""
    },
    "op_rw_latency_out_bytes_histogram": {
        "type": 18,
        "description": "Histogram of rw operation latency (including queue time)  +  da
        "nick": ""
    }
}
}

```

DUMP

The actual dump is similar to the schema, except that there are actual value groups. For example:

```

"osd": {
  "op_r_latency_out_bytes_histogram": {
    "axes": [
      {
        "name": "Latency (usec)",
        "min": 0,
        "quant_size": 100000,
        "buckets": 32,
        "scale_type": "log2",
        "ranges": [
          {
            "max": -1
          },
          {
            "min": 0,
            "max": 99999
          },
          {
            "min": 100000,
            "max": 199999
          },
          {
            "min": 200000,
            "max": 399999
          },
          {
            "min": 400000,
            "max": 799999
          },
          {
            "min": 800000,
            "max": 1599999
          },
          {
            "min": 1600000,
            "max": 3199999
          },
          {
            "min": 3200000,
            "max": 6399999
          },
          {
            "min": 6400000,
            "max": 12799999
          },
          {
            "min": 12800000,
            "max": 25599999
          },
          {
            "min": 25600000,
            "max": 51199999
          },
          {
            "min": 51200000,
            "max": 102399999
          },
          {

```

```
{
  "min": 102400000,
  "max": 204799999
},
{
  "min": 204800000,
  "max": 409599999
},
{
  "min": 409600000,
  "max": 819199999
},
{
  "min": 819200000,
  "max": 1638399999
},
{
  "min": 1638400000,
  "max": 3276799999
},
{
  "min": 3276800000,
  "max": 6553599999
},
{
  "min": 6553600000,
  "max": 13107199999
},
{
  "min": 13107200000,
  "max": 26214399999
},
{
  "min": 26214400000,
  "max": 52428799999
},
{
  "min": 52428800000,
  "max": 104857599999
},
{
  "min": 104857600000,
  "max": 209715199999
},
{
  "min": 209715200000,
  "max": 419430399999
},
{
  "min": 419430400000,
  "max": 838860799999
},
{
  "min": 838860800000,
  "max": 1677721599999
},
{
  "min": 1677721600000,
  "max": 3355443199999
},
{
  "min": 3355443200000,
  "max": 6710886399999
},
{
  "min": 6710886400000,
  "max": 13421772799999
},
{
  "min": 13421772800000,
  "max": 26843545599999
},
{
  "min": 26843545600000,
  "max": 53687091199999
},
},
```

```

    },
    {
      "min": 53687091200000
    }
  ]
},
{
  "name": "Request size (bytes)",
  "min": 0,
  "quant_size": 512,
  "buckets": 32,
  "scale_type": "log2",
  "ranges": [
    {
      "max": -1
    },
    {
      "min": 0,
      "max": 511
    },
    {
      "min": 512,
      "max": 1023
    },
    {
      "min": 1024,
      "max": 2047
    },
    {
      "min": 2048,
      "max": 4095
    },
    {
      "min": 4096,
      "max": 8191
    },
    {
      "min": 8192,
      "max": 16383
    },
    {
      "min": 16384,
      "max": 32767
    },
    {
      "min": 32768,
      "max": 65535
    },
    {
      "min": 65536,
      "max": 131071
    },
    {
      "min": 131072,
      "max": 262143
    },
    {
      "min": 262144,
      "max": 524287
    },
    {
      "min": 524288,
      "max": 1048575
    },
    {
      "min": 1048576,
      "max": 2097151
    },
    {
      "min": 2097152,
      "max": 4194303
    },
    {
      "min": 4194304,
      "max": 8388607
    },
  ],

```

```
{
    {
        "min": 8388608,
        "max": 16777215
    },
    {
        "min": 16777216,
        "max": 33554431
    },
    {
        "min": 33554432,
        "max": 67108863
    },
    {
        "min": 67108864,
        "max": 134217727
    },
    {
        "min": 134217728,
        "max": 268435455
    },
    {
        "min": 268435456,
        "max": 536870911
    },
    {
        "min": 536870912,
        "max": 1073741823
    },
    {
        "min": 1073741824,
        "max": 2147483647
    },
    {
        "min": 2147483648,
        "max": 4294967295
    },
    {
        "min": 4294967296,
        "max": 8589934591
    },
    {
        "min": 8589934592,
        "max": 17179869183
    },
    {
        "min": 17179869184,
        "max": 34359738367
    },
    {
        "min": 34359738368,
        "max": 68719476735
    },
    {
        "min": 68719476736,
        "max": 137438953471
    },
    {
        "min": 137438953472,
        "max": 274877906943
    },
    {
        "min": 274877906944
    }
}

],
"values": [
    0,
    0,
    0,
    0,
    0,
    0,
    0,
    0,
```


[illegible]

[illegible]

[illegible]

This represents the 2d histogram, consisting of 9 history entires and 32 value groups per each history entry. "Ranges"

element denote value bounds for each of value groups. “Buckets” denote amount of value groups (“buckets”), “Min” is a minimum accepted valaue, “quant_size” is quantization unit and “scale_type” is either “log2” (logarhitmic scale) or “linear” (linear scale). You can use histogram_dump.py tool (see src/tools/histogram_dump.py) for quick visualisation of existing histogram data.
