

PUBLIC OSD VERSION

We maintain two versions on disk: an `eversion_t pg_log.head` and a `version_t info.user_version`. Each object is tagged with both the pg version and user_version it was last modified with. The PG version is modified by manipulating `OpContext::at_version` and then persisting it to the pg log as transactions, and is incremented in all the places it used to be. The user_version is modified by manipulating the new `OpContext::user_at_version` and is also persisted via the pg log transactions. `user_at_version` is modified only in `PrimaryLogPG::prepare_transaction` when the op was a “user modify” (a non-watch write), and the durable user_version is updated according to the following rules: 1) set `user_at_version` to the maximum of `ctx->new_obs.oi.user_version+1` and `info.last_user_version+1`. 2) set `user_at_version` to the maximum of itself and `ctx->at_version.version`. 3) `ctx->new_obs.oi.user_version = ctx->user_at_version` (to change the object’s user_version)

This set of update semantics mean that for traditional pools the user_version will be equal to the past `reassert_version`, while for caching pools the object and PG user-version will be able to cross pools without making a total mess of things. In order to support old clients, we keep the old `reassert_version` but rename it to “`bad_replay_version`”; we fill it in as before: for writes it is set to the `at_version` (and is the proper replay version); for watches it is set to our user version; for ENOENT replies it is set to the replay version’s epoch but the user_version’s version. We also now fill in the `version_t` portion of the `bad_replay_version` on read ops as well as write ops, which should be fine for all old clients.

For new clients, we prevent them from reading `bad_replay_version` and add two proper members: `user_version` and `replay_version`; `user_version` is filled in on every operation (reads included) while `replay_version` is filled in for writes.

The `objclass` function `get_current_version()` now always returns the `pg->info.last_user_version`, which means it is guaranteed to contain the version of the last user update in the PG (including on reads!).