

# CEPH-OSD – CEPH OBJECT STORAGE DAEMON

## SYNOPSIS

**ceph-osd** -i *osdnum* [ -osd-data *datapath* ] [ -osd-journal *journal* ] [ -mkfs ] [ -mkjournal ] [-flush-journal] [-check-allows-journal] [-check-wants-journal] [-check-needs-journal] [ -mkkey ]

## DESCRIPTION

**ceph-osd** is the object storage daemon for the Ceph distributed file system. It is responsible for storing objects on a local file system and providing access to them over the network.

The *datapath* argument should be a directory on a xfs file system where the object data resides. The *journal* is optional, and is only useful performance-wise when it resides on a different disk than *datapath* with low latency (ideally, an NVRAM device).

## OPTIONS

### -f, --foreground

Foreground: do not daemonize after startup (run in foreground). Do not generate a pid file. Useful when run via **ceph-run(8)**.

### -d

Debug mode: like -f, but also send all log output to stderr.

### --setuser userorgid

Set uid after starting. If a username is specified, the user record is looked up to get a uid and a gid, and the gid is also set as well, unless -setgroup is also specified.

### --setgroup grouporgid

Set gid after starting. If a group name is specified the group record is looked up to get a gid.

### --osd-data osddata

Use object store at *osddata*.

### --osd-journal journal

Journal updates to *journal*.

### --check-wants-journal

Check whether a journal is desired.

### --check-allows-journal

Check whether a journal is allowed.

### --check-needs-journal

Check whether a journal is required.

### --mkfs

Create an empty object repository. This also initializes the journal (if one is defined).

### --mkkey

Generate a new secret key. This is normally used in combination with - -mkfs as it is more convenient than generating a key by hand with **ceph-authtool(8)**.

### --mkjournal

Create a new journal file to match an existing object repository. This is useful if the journal device or file is wiped out due to a disk or file system failure.

### --flush-journal

Flush the journal to permanent store. This runs in the foreground so you know when it's completed. This can be useful if you want to resize the journal or need to otherwise destroy it: this guarantees you won't lose data.

## **--get-cluster-fsid**

Print the cluster fsid (uuid) and exit.

## **--get-osd-fsid**

Print the OSD's fsid and exit. The OSD's uuid is generated at `-mkfs` time and is thus unique to a particular instantiation of this OSD.

## **--get-journal-fsid**

Print the journal's uuid. The journal fsid is set to match the OSD fsid at `-mkfs` time.

## **-c** `ceph.conf`, **--conf**=`ceph.conf`

Use *ceph.conf* configuration file instead of the default `/etc/ceph/ceph.conf` for runtime configuration options.

## **-m** `monaddress[:port]`

Connect to specified monitor (instead of looking through `ceph.conf`).

## AVAILABILITY

**ceph-osd** is part of Ceph, a massively scalable, open-source, distributed storage system. Please refer to the Ceph documentation at <http://ceph.com/docs> for more information.

## SEE ALSO

[ceph\(8\)](#), [ceph-mds\(8\)](#), [ceph-mon\(8\)](#), [ceph-authtool\(8\)](#)

---