

MOUNT.CEPH – MOUNT A CEPH FILE SYSTEM

SYNOPSIS

mount.ceph *monaddr1[,monaddr2,...]:/[subdir] dir [-o options]*

DESCRIPTION

mount.ceph is a simple helper for mounting the Ceph file system on a Linux host. It serves to resolve monitor hostname(s) into IP addresses and read authentication keys from disk; the Linux kernel client component does most of the real work. In fact, it is possible to mount a non-authenticated Ceph file system without **mount.ceph** by specifying monitor address(es) by IP:

```
mount -t ceph 1.2.3.4:/ mountpoint
```

Each monitor address *monaddr* takes the form *host[:port]*. If the port is not specified, the Ceph default of 6789 is assumed.

Multiple monitor addresses can be separated by commas. Only one responsible monitor is needed to successfully mount; the client will learn about all monitors from any responsive monitor. However, it is a good idea to specify more than one in case one happens to be down at the time of mount.

A subdirectory *subdir* may be specified if a subset of the file system is to be mounted.

Mount helper application conventions dictate that the first two options are device to be mounted and destination path. Options must be passed only after these fixed arguments.

OPTIONS

wsiz

int (bytes), max write size. Default: none (writeback uses smaller of *wsiz* and stripe unit)

rsiz

int (bytes), max read size. Default: none

rasiz

int (bytes), max readahead, multiple of 1024, Default: 8388608 (8192*1024)

osdtimeout

int (seconds), Default: 60

osdkeepalive

int, Default: 5

mount_timeout

int (seconds), Default: 60

osd_idle_ttl

int (seconds), Default: 60

caps_wanted_delay_min

int, cap release delay, Default: 5

caps_wanted_delay_max

int, cap release delay, Default: 60

cap_release_safety

int, Default: calculated

readdir_max_entries

int, Default: 1024

readdir_max_bytes

int, Default: 524288 (512*1024)

write_congestion_kb

int (kb), max writeback in flight. scale with available memory. Default: calculated from available memory

snapdirname

string, set the name of the hidden snapdir. Default: .snap

name

RADOS user to authenticate as when using cephx. Default: guest

secret

secret key for use with cephx. This option is insecure because it exposes the secret on the command line. To avoid this, use the secretfile option.

secretfile

path to file containing the secret key to use with cephx

ip

my ip

noshare

create a new client instance, instead of sharing an existing instance of a client mounting the same cluster

dirstat

funky *cat dirname* for stats, Default: off

nodirstat

no funky *cat dirname* for stats

rbytes

Report the recursive size of the directory contents for st_size on directories. Default: on

norbytes

Do not report the recursive size of the directory contents for st_size on directories.

nocrc

no data crc on writes

noasyncreaddir

no dcache readdir

EXAMPLES

Mount the full file system:

```
mount.ceph monhost:/ /mnt/foo
```

If there are multiple monitors:

```
mount.ceph monhost1,monhost2,monhost3:/ /mnt/foo
```

If **ceph-mon**(8) is running on a non-standard port:

```
mount.ceph monhost1:7000,monhost2:7000,monhost3:7000:/ /mnt/foo
```

To mount only part of the namespace:

```
mount.ceph monhost1:/some/small/thing /mnt/thing
```

Assuming mount.ceph(8) is installed properly, it should be automatically invoked by mount(8) like so:

```
mount -t ceph monhost:/ /mnt/foo
```

AVAILABILITY

mount.ceph is part of Ceph, a massively scalable, open-source, distributed storage system. Please refer to the Ceph

documentation at <http://ceph.com/docs> for more information.

SEE ALSO

[ceph-fuse\(8\)](#), [ceph\(8\)](#)