

OSDMAPTOOL – CEPH OSD CLUSTER MAP MANIPULATION TOOL

SYNOPSIS

osdmapprool *mapfilename* [-print] [-createsimple *numosd* [-pgbits *bitsperosd*]] [-clobber]

DESCRIPTION

osdmapprool is a utility that lets you create, view, and manipulate OSD cluster maps from the Ceph distributed storage system. Notably, it lets you extract the embedded CRUSH map or import a new CRUSH map.

OPTIONS

--print

will simply make the tool print a plaintext dump of the map, after any modifications are made.

--clobber

will allow osdmapprool to overwrite mapfilename if changes are made.

--import-crush mapfile

will load the CRUSH map from mapfile and embed it in the OSD map.

--export-crush mapfile

will extract the CRUSH map from the OSD map and write it to mapfile.

--createsimple numosd [--pgbits bitsperosd]

will create a relatively generic OSD map with the numosd devices. If -pgbits is specified, the initial placement group counts will be set with bitsperosd bits per OSD. That is, the pg_num map attribute will be set to numosd shifted by bitsperosd.

--test-map-pgs [--pool poolid]

will print out the mappings from placement groups to OSDs.

--test-map-pgs-dump [--pool poolid]

will print out the summary of all placement groups and the mappings from them to the mapped OSDs.

EXAMPLE

To create a simple map with 16 devices:

```
osdmapprool --createsimple 16 osdmap --clobber
```

To view the result:

```
osdmapprool --print osdmap
```

To view the mappings of placement groups for pool 0:

```
osdmapprool --test-map-pgs-dump rbd --pool 0
```

```
pool 0 pg_num 8
0.0    [0,2,1] 0
0.1    [2,0,1] 2
0.2    [0,1,2] 0
0.3    [2,0,1] 2
0.4    [0,2,1] 0
0.5    [0,2,1] 0
0.6    [0,1,2] 0
```

```
0.7 [1,0,2] 1
#osd count first primary c wt wt
osd.0 8 5 5 1 1
osd.1 8 1 1 1 1
osd.2 8 2 2 1 1
in 3
avg 8 stddev 0 (0x) (expected 2.3094 0.288675x)
min osd.0 8
max osd.0 8
size 0 0
size 1 0
size 2 0
size 3 8
```

In which,

1. pool 0 has 8 placement groups. And two tables follow:
2. A table for placement groups. Each row presents a placement group. With columns of:
 - placement group id,
 - acting set, and
 - primary OSD.
3. A table for all OSDs. Each row presents an OSD. With columns of:
 - count of placement groups being mapped to this OSD,
 - count of placement groups where this OSD is the first one in their acting sets,
 - count of placement groups where this OSD is the primary of them,
 - the CRUSH weight of this OSD, and
 - the weight of this OSD.
4. Looking at the number of placement groups held by 3 OSDs. We have
 - avarge, stddev, stddev/average, expected stddev, expected stddev / average
 - min and max
5. The number of placement groups mapping to n OSDs. In this case, all 8 placement groups are mapping to 3 different OSDs.

In a less-balanced cluster, we could have following output for the statistics of placement group distribution, whose standard deviation is 1.41421:

```
#osd count first primary c wt wt
osd.0 8 5 5 1 1
osd.1 8 1 1 1 1
osd.2 8 2 2 1 1

#osd count first primary c wt wt
osd.0 33 9 9 0.0145874 1
osd.1 34 14 14 0.0145874 1
osd.2 31 7 7 0.0145874 1
osd.3 31 13 13 0.0145874 1
osd.4 30 14 14 0.0145874 1
osd.5 33 7 7 0.0145874 1
in 6
avg 32 stddev 1.41421 (0.0441942x) (expected 5.16398 0.161374x)
min osd.4 30
max osd.1 34
size 00
size 10
size 20
size 364
```

AVAILABILITY

osdmapprootool is part of Ceph, a massively scalable, open-source, distributed storage system. Please refer to the Ceph documentation at <http://ceph.com/docs> for more information.

SEE ALSO

ceph(8), **crushtool(8)**,

