

RADOSGW-ADMIN – RADOS REST GATEWAY USER ADMINISTRATION UTILITY

SYNOPSIS

radosgw-admin *command* [*options ...*]

DESCRIPTION

radosgw-admin is a RADOS gateway user administration utility. It allows creating and modifying users.

COMMANDS

radosgw-admin utility uses many commands for administration purpose which are as follows:

user create

Create a new user.

user modify

Modify a user.

user info

Display information of a user, and any potentially available subusers and keys.

user rm

Remove a user.

user suspend

Suspend a user.

user enable

Re-enable user after suspension.

user check

Check user info.

user stats

Show user stats as accounted by quota subsystem.

user list

List all users.

caps add

Add user capabilities.

caps rm

Remove user capabilities.

subuser create

Create a new subuser (primarily useful for clients using the Swift API).

subuser modify

Modify a subuser.

subuser rm

Remove a subuser.

key create

Create access key.

key rm

Remove access key.

bucket list

List all buckets.

bucket limit check

Show bucket sharding stats.

bucket link

Link bucket to specified user.

bucket unlink

Unlink bucket from specified user.

bucket stats

Returns bucket statistics.

bucket rm

Remove a bucket.

bucket check

Check bucket index.

bucket rewrite

Rewrite all objects in the specified bucket.

bucket reshard

Reshard a bucket.

bucket sync disable

Disable bucket sync.

bucket sync enable

Enable bucket sync.

bi get

Retrieve bucket index object entries.

bi put

Store bucket index object entries.

bi list

List raw bucket index entries.

bi purge

Purge bucket index entries.

object rm

Remove an object.

object stat

Stat an object for its metadata.

object unlink

Unlink object from bucket index.

object rewrite

Rewrite the specified object.

objects expire

Run expired objects cleanup.

period delete

Delete a period.

period get

Get the period info.

period get-current

Get the current period info.

period pull

Pull a period.

period push

Push a period.

period list

List all periods.

period update

Update the staging period.

period commit

Commit the staging period.

quota set

Set quota params.

quota enable

Enable quota.

quota disable

Disable quota.

global quota get

View global quota parameters.

global quota set

Set global quota parameters.

global quota enable

Enable a global quota.

global quota disable

Disable a global quota.

realm create

Create a new realm.

realm delete

Delete a realm.

realm get

Show the realm info.

realm get-default

Get the default realm name.

realm list

List all realms.

realm list-periods

List all realm periods.

realm rename

Rename a realm.

realm set

Set the realm info (requires infile).

realm default

Set the realm as default.

realm pull

Pull a realm and its current period.

zonegroup add

Add a zone to a zonegroup.

zonegroup create

Create a new zone group info.

zonegroup default

Set the default zone group.

zonegroup delete

Delete a zone group info.

zonegroup get

Show the zone group info.

zonegroup modify

Modify an existing zonegroup.

zonegroup set

Set the zone group info (requires infile).

zonegroup remove

Remove a zone from a zonegroup.

zonegroup rename

Rename a zone group.

zonegroup list

List all zone groups set on this cluster.

zonegroup placement list

List zonegroup's placement targets.

zonegroup placement add

Add a placement target id to a zonegroup.

zonegroup placement modify

Modify a placement target of a specific zonegroup.

zonegroup placement rm

Remove a placement target from a zonegroup.

zonegroup placement default

Set a zonegroup's default placement target.

zone create

Create a new zone.

zone delete

Delete a zone.

zone get

Show zone cluster params.

zone set

Set zone cluster params (requires infile).

zone modify

Modify an existing zone.

zone list

List all zones set on this cluster.

metadata sync status

Get metadata sync status.

metadata sync init

Init metadata sync.

metadata sync run

Run metadata sync.

data sync status

Get data sync status of the specified source zone.

data sync init

Init data sync for the specified source zone.

data sync run

Run data sync for the specified source zone.

sync error list

list sync error.

sync error trim

trim sync error.

zone rename

Rename a zone.

zone placement list

List zone's placement targets.

zone placement add

Add a zone placement target.

zone placement modify

Modify a zone placement target.

zone placement rm

Remove a zone placement target.

pool add

Add an existing pool for data placement.

pool rm

Remove an existing pool from data placement set.

pools list

List placement active set.

policy

Display bucket/object policy.

log list

List log objects.

log show

Dump a log from specific object or (bucket + date + bucket-id). (NOTE: required to specify formatting of date to “YYYY-MM-DD-hh”)

log rm

Remove log object.

usage show

Show the usage information (with optional user and date range).

usage trim

Trim usage information (with optional user and date range).

gc list

Dump expired garbage collection objects (specify `-include-all` to list all entries, including unexpired).

gc process

Manually process garbage.

lc list

List all bucket lifecycle progress.

lc process

Manually process lifecycle.

metadata get

Get metadata info.

metadata put

Put metadata info.

metadata rm

Remove metadata info.

metadata list

List metadata info.

mdlog list

List metadata log.

mdlog trim

Trim metadata log.

mdlog status

Read metadata log status.

bilog list

List bucket index log.

bilog trim

Trim bucket index log (use start-marker, end-marker).

datalog list

List data log.

datalog trim

Trim data log.

datalog status

Read data log status.

opstate list

List stateful operations entries (use client_id, op_id, object).

opstate set

Set state on an entry (use client_id, op_id, object, state).

opstate renew

Renew state on an entry (use client_id, op_id, object).

opstate rm

Remove entry (use client_id, op_id, object).

replicalog get

Get replica metadata log entry.

replicalog update

Update replica metadata log entry.

replicalog delete

Delete replica metadata log entry.

orphans find

Init and run search for leaked rados objects

orphans finish

Clean up search for leaked rados objects

orphans list-jobs

List the current job-ids for the orphans search.

role create

create a new AWS role for use with STS.

role delete

Delete a role.

role get

Get a role.

role list

List the roles with specified path prefix.

role modify

Modify the assume role policy of an existing role.

role-policy put

Add/update permission policy to role.

role-policy list

List the policies attached to a role.

role-policy get

Get the specified inline policy document embedded with the given role.

role-policy delete

Delete the policy attached to a role

reshard add

Schedule a resharding of a bucket

reshard list

List all bucket resharding or scheduled to be resharded

reshard process

Process of scheduled reshard jobs

reshard status

Resharding status of a bucket

reshard cancel

Cancel resharding a bucket

OPTIONS

-c ceph.conf, **--conf**=ceph.conf

Use ceph.conf configuration file instead of the default /etc/ceph/ceph.conf to determine monitor addresses during startup.

-m monaddress[:port]
Connect to specified monitor (instead of looking through ceph.conf).

--tenant=<tenant>
Name of the tenant.

--uid=uid
The radosgw user ID.

--subuser=<name>
Name of the subuser.

--access-key=<key>
S3 access key.

--email=email
The e-mail address of the user.

--secret/--secret-key=<key>
The secret key.

--gen-access-key
Generate random access key (for S3).

--gen-secret
Generate random secret key.

--key-type=<type>
key type, options are: swift, s3.

--temp-url-key[-2]=<key>
Temporary url key.

--max-buckets
max number of buckets for a user (0 for no limit, negative value to disable bucket creation). Default is 1000.

--access=<access>
Set the access permissions for the sub-user. Available access permissions are read, write, readwrite and full.

--display-name=<name>
The display name of the user.

--admin
Set the admin flag on the user.

--system
Set the system flag on the user.

--bucket=bucket
Specify the bucket name.

--pool=<pool>
Specify the pool name. Also used with *orphans find* as data pool to scan for leaked rados objects.

--object=object
Specify the object name.

--date=yyyy-mm-dd
The date in the format yyyy-mm-dd.

--start-date=yyyy-mm-dd
The start date in the format yyyy-mm-dd.

--end-date=yyyy-mm-dd
The end date in the format yyyy-mm-dd.

--bucket-id=<bucket-id>

Specify the bucket id.

--shard-id=<shard-id>

Optional for mdlog list. Required for mdlog trim, replica mdlog get/delete, replica datalog get/delete.

--auth-uid=auid

The librados auid.

--purge-data

When specified, user removal will also purge all the user data.

--purge-keys

When specified, subuser removal will also purge all the subuser keys.

--purge-objects

When specified, the bucket removal will also purge all objects in it.

--metadata-key=<key>

Key to retrieve metadata from with metadata get.

--remote=<remote>

Zone or zonegroup id of remote gateway.

--period=<id>

Period id.

--epoch=<number>

Period epoch.

--commit

Commit the period during 'period update'.

--staging

Get the staging period info.

--master

Set as master.

--master-zone=<id>

Master zone id.

--rgw-realm=<name>

The realm name.

--realm-id=<id>

The realm id.

--realm-new-name=<name>

New name of realm.

--rgw-zonegroup=<name>

The zonegroup name.

--zonegroup-id=<id>

The zonegroup id.

--zonegroup-new-name=<name>

The new name of the zonegroup.

--rgw-zone=<zone>

Zone in which radosgw is running.

--zone-id=<id>

The zone id.

--zone-new-name=<name>
The new name of the zone.

--source-zone
The source zone for data sync.

--default
Set the entity (realm, zonegroup, zone) as default.

--read-only
Set the zone as read-only when adding to the zonegroup.

--placement-id
Placement id for the zonegroup placement commands.

--tags=<list>
The list of tags for zonegroup placement add and modify commands.

--tags-add=<list>
The list of tags to add for zonegroup placement modify command.

--tags-rm=<list>
The list of tags to remove for zonegroup placement modify command.

--endpoints=<list>
The zone endpoints.

--index-pool=<pool>
The placement target index pool.

--data-pool=<pool>
The placement target data pool.

--data-extra-pool=<pool>
The placement target data extra (non-ec) pool.

--placement-index-type=<type>
The placement target index type (normal, indexless, or #id).

--tier-type=<type>
The zone tier type.

--tier-config=<k>=<v>[, ...]
Set zone tier config keys, values.

--tier-config-rm=<k>[, ...]
Unset zone tier config keys.

--sync-from-all[=false]
Set/reset whether zone syncs from all zonegroup peers.

--sync-from=[zone-name] [, ...]
Set the list of zones to sync from.

--sync-from-rm=[zone-name] [, ...]
Remove the zones from list of zones to sync from.

--fix
Besides checking bucket index, will also fix it.

--check-objects
bucket check: Rebuilds bucket index according to actual objects state.

--format=<format>
Specify output format for certain operations. Supported formats: xml, json.

--sync-stats

Option for ‘user stats’ command. When specified, it will update user stats with the current stats reported by user’s buckets indexes.

--show-log-entries=<flag>

Enable/disable dump of log entries on log show.

--show-log-sum=<flag>

Enable/disable dump of log summation on log show.

--skip-zero-entries

Log show only dumps entries that don’t have zero value in one of the numeric field.

--infile

Specify a file to read in when setting data.

--state=<state>

Specify a state for the opstate set command.

--replica-log-type <logtypestr>

Replica log type (metadata, data, bucket), required for replica log operations.

--categories=<list>

Comma separated list of categories, used in usage show.

--caps=<caps>

List of caps (e.g., “usage=read, write; user=read”).

--compression=<compression-algorithm>

Placement target compression algorithm (lz4|snappy|zlib|zstd)

--yes-i-really-mean-it

Required for certain operations.

--min-rewrite-size

Specify the min object size for bucket rewrite (default 4M).

--max-rewrite-size

Specify the max object size for bucket rewrite (default ULLONG_MAX).

--min-rewrite-stripe-size

Specify the min stripe size for object rewrite (default 0). If the value is set to 0, then the specified object will always be rewritten for restriping.

--warnings-only

When specified with bucket limit check, list only buckets nearing or over the current max objects per shard value.

--bypass-gc

When specified with bucket deletion, triggers object deletions by not involving GC.

--inconsistent-index

When specified with bucket deletion and bypass-gc set to true, ignores bucket index consistency.

QUOTA OPTIONS

--max-objects

Specify max objects (negative value to disable).

--max-size

Specify max size (in B/K/M/G/T, negative value to disable).

--quota-scope

The scope of quota (bucket, user).

ORPHANS SEARCH OPTIONS

--num-shards

Number of shards to use for keeping the temporary scan info

--orphan-stale-secs

Number of seconds to wait before declaring an object to be an orphan. Default is 86400 (24 hours).

--job-id

Set the job id (for orphans find)

--max-concurrent-ios

Maximum concurrent ios for orphans find. Default is 32.

ORPHANS LIST-JOBS OPTIONS

--extra-info

Provide extra info in the job list.

ROLE OPTIONS

--role-name

The name of the role to create.

--path

The path to the role.

--assume-role-policy-doc

The trust relationship policy document that grants an entity permission to assume the role.

--policy-name

The name of the policy document.

--policy-doc

The permission policy document.

--path-prefix

The path prefix for filtering the roles.

EXAMPLES

Generate a new user:

```
$ radosgw-admin user create --display-name="johnny rotten" --uid=johnny
{ "user_id": "johnny",
  "rados_uid": 0,
  "display_name": "johnny rotten",
  "email": "",
  "suspended": 0,
  "subusers": [],
  "keys": [
    { "user": "johnny",
      "access_key": "TCICW53D9BQ2VGC46I44",
      "secret_key": "tfm9aHMI8X76L3UdgE+ZQaJag1vJQmE6HDb5Lbrz"}],
  "swift_keys": []}
```

Remove a user:

```
$ radosgw-admin user rm --uid=johnny
```

Remove a user and all associated buckets with their contents:

```
$ radosgw-admin user rm --uid=johnny --purge-data
```

Remove a bucket:

```
$ radosgw-admin bucket rm --bucket=foo
```

Link bucket to specified user:

```
$ radosgw-admin bucket link --bucket=foo --bucket_id=<bucket id> --uid=johnny
```

Unlink bucket from specified user:

```
$ radosgw-admin bucket unlink --bucket=foo --uid=johnny
```

Show the logs of a bucket from April 1st, 2012:

```
$ radosgw-admin log show --bucket=foo --date=2012-04-01-01 --bucket-id=default.14193.1
```

Show usage information for user from March 1st to (but not including) April 1st, 2012:

```
$ radosgw-admin usage show --uid=johnny \  
    --start-date=2012-03-01 --end-date=2012-04-01
```

Show only summary of usage information for all users:

```
$ radosgw-admin usage show --show-log-entries=false
```

Trim usage information for user until March 1st, 2012:

```
$ radosgw-admin usage trim --uid=johnny --end-date=2012-04-01
```

AVAILABILITY

radosgw-admin is part of Ceph, a massively scalable, open-source, distributed storage system. Please refer to the Ceph documentation at <http://ceph.com/docs> for more information.

SEE ALSO

[ceph\(8\)](#) [radosgw\(8\)](#)