

MONITORING THE ISCSI GATEWAYS

Ceph provides an additional tool for iSCSI gateway environments to monitor performance of exported RADOS Block Device (RBD) images.

The gwtop tool is a top-like tool that displays aggregated performance metrics of RBD images that are exported to clients over iSCSI. The metrics are sourced from a Performance Metrics Domain Agent (PMDA). Information from the Linux-IO target (LIO) PMDA is used to list each exported RBD image with the connected client and its associated I/O metrics.

Requirements:

- A running Ceph iSCSI gateway

Installing:

1. As root, install the ceph-iscsi-tools package on each iSCSI gateway node:

```
# yum install ceph-iscsi-tools
```

2. As root, install the performance co-pilot package on each iSCSI gateway node:

```
# yum install pcp
```

3. As root, install the LIO PMDA package on each iSCSI gateway node:

```
# yum install pcp-pmda-lio
```

4. As root, enable and start the performance co-pilot service on each iSCSI gateway node:

```
# systemctl enable pmcd
# systemctl start pmcd
```

5. As root, register the pcp-pmda-lio agent:

```
cd /var/lib/pcp/pmdas/lio
./Install
```

By default, gwtop assumes the iSCSI gateway configuration object is stored in a RADOS object called gateway.conf in the rbd pool. This configuration defines the iSCSI gateways to contact for gathering the performance statistics. This can be overridden by using either the -g or -c flags. See gwtop --help for more details.

The LIO configuration determines which type of performance statistics to extract from performance co-pilot. When gwtop starts it looks at the LIO configuration, and if it find user-space disks, then gwtop selects the LIO collector automatically.

Example ``gwtop`` Outputs

```
gwtop  2/2 Gateways  CPU% MIN:  4 MAX:  5  Network Total In:   2M Out:   3M 10:20:00
Capacity: 8G      Disks:  8  IOPS: 503  Clients:  1 Ceph: HEALTH_OK      OSDs:  3
Pool.Image      Src      Size      iops      rMB/s      wMB/s      Client
iscsi.t1703                500M        0      0.00      0.00
iscsi.testme1          500M        0      0.00      0.00
iscsi.testme2          500M        0      0.00      0.00
iscsi.testme3          500M        0      0.00      0.00
iscsi.testme5          500M        0      0.00      0.00
rbd.myhost_1      T        4G      504      1.95      0.00      rh460p (CON)
rbd.test_2                1G         0      0.00      0.00
rbd.testme          500M        0      0.00      0.00
```

In the *Client* column, (CON) means the iSCSI initiator (client) is currently logged into the iSCSI gateway. If -multi- is displayed, then multiple clients are mapped to the single RBD image.

