


BLOCK DEVICES AND CLOUDSTACK

You may use Ceph Block Device images with CloudStack 4.0 and higher through `libvirt`, which configures the QEMU interface to `librbd`. Ceph stripes block device images as objects across the cluster, which means that large Ceph Block Device images have better performance than a standalone server!

To use Ceph Block Devices with CloudStack 4.0 and higher, you must install QEMU, `libvirt`, and CloudStack first. We recommend using a separate physical host for your CloudStack installation. CloudStack recommends a minimum of 4GB of RAM and a dual-core processor, but more CPU and RAM will perform better. The following diagram depicts the CloudStack/Ceph technology stack.



Important: To use Ceph Block Devices with CloudStack, you must have access to a running Ceph Storage Cluster.

CloudStack integrates with Ceph's block devices to provide CloudStack with a back end for CloudStack's Primary Storage. The instructions below detail the setup for CloudStack Primary Storage.

Note: We recommend installing with Ubuntu 14.04 or later so that you can use package installation instead of having to compile `libvirt` from source.

Installing and configuring QEMU for use with CloudStack doesn't require any special handling. Ensure that you have a running Ceph Storage Cluster. Install QEMU and configure it for use with Ceph; then, install `libvirt` version 0.9.13 or higher (you may need to compile from source) and ensure it is running with Ceph.

Note: Ubuntu 14.04 and CentOS 7.2 will have `libvirt` with RBD storage pool support enabled by default.

CREATE A POOL

By default, Ceph block devices use the `rbd` pool. Create a pool for CloudStack NFS Primary Storage. Ensure your Ceph cluster is running, then create the pool.

```
ceph osd pool create cloudstack
```

See [Create a Pool](#) for details on specifying the number of placement groups for your pools, and [Placement Groups](#) for details on the number of placement groups you should set for your pools.

A newly created pool must be initialized prior to use. Use the `rbd` tool to initialize the pool:

```
rbd pool init cloudstack
```

CREATE A CEPH USER

To access the Ceph cluster we require a Ceph user which has the correct credentials to access the `cloudstack` pool we just created. Although we could use `client.admin` for this, it's recommended to create a user with only access to the `cloudstack` pool.

```
ceph auth get-or-create client.cloudstack mon 'profile rbd' osd 'profile rbd pool=cloudstack'
```

Use the information returned by the command in the next step when adding the Primary Storage.

See [User Management](#) for additional details.

ADD PRIMARY STORAGE

To add primary storage, refer to [Add Primary Storage \(4.2.0\)](#) to add a Ceph block device, the steps include:

1. Log in to the CloudStack UI.
2. Click **Infrastructure** on the left side navigation bar.
3. Select the Zone you want to use for Primary Storage.
4. Click the **Compute** tab.
5. Select **View All** on the *Primary Storage* node in the diagram.
6. Click **Add Primary Storage**.
7. Follow the CloudStack instructions.
 - For **Protocol**, select RBD.
 - Add cluster information (cephx is supported). Note: Do not include the `client.` part of the user.
 - Add rbd as a tag.

CREATE A DISK OFFERING

To create a new disk offering, refer to [Create a New Disk Offering \(4.2.0\)](#). Create a disk offering so that it matches the rbd tag. The StoragePoolAllocator will choose the rbd pool when searching for a suitable storage pool. If the disk offering doesn't match the rbd tag, the StoragePoolAllocator may select the pool you created (e.g., cloudstack).

LIMITATIONS

- CloudStack will only bind to one monitor (You can however create a Round Robin DNS record over multiple monitors)